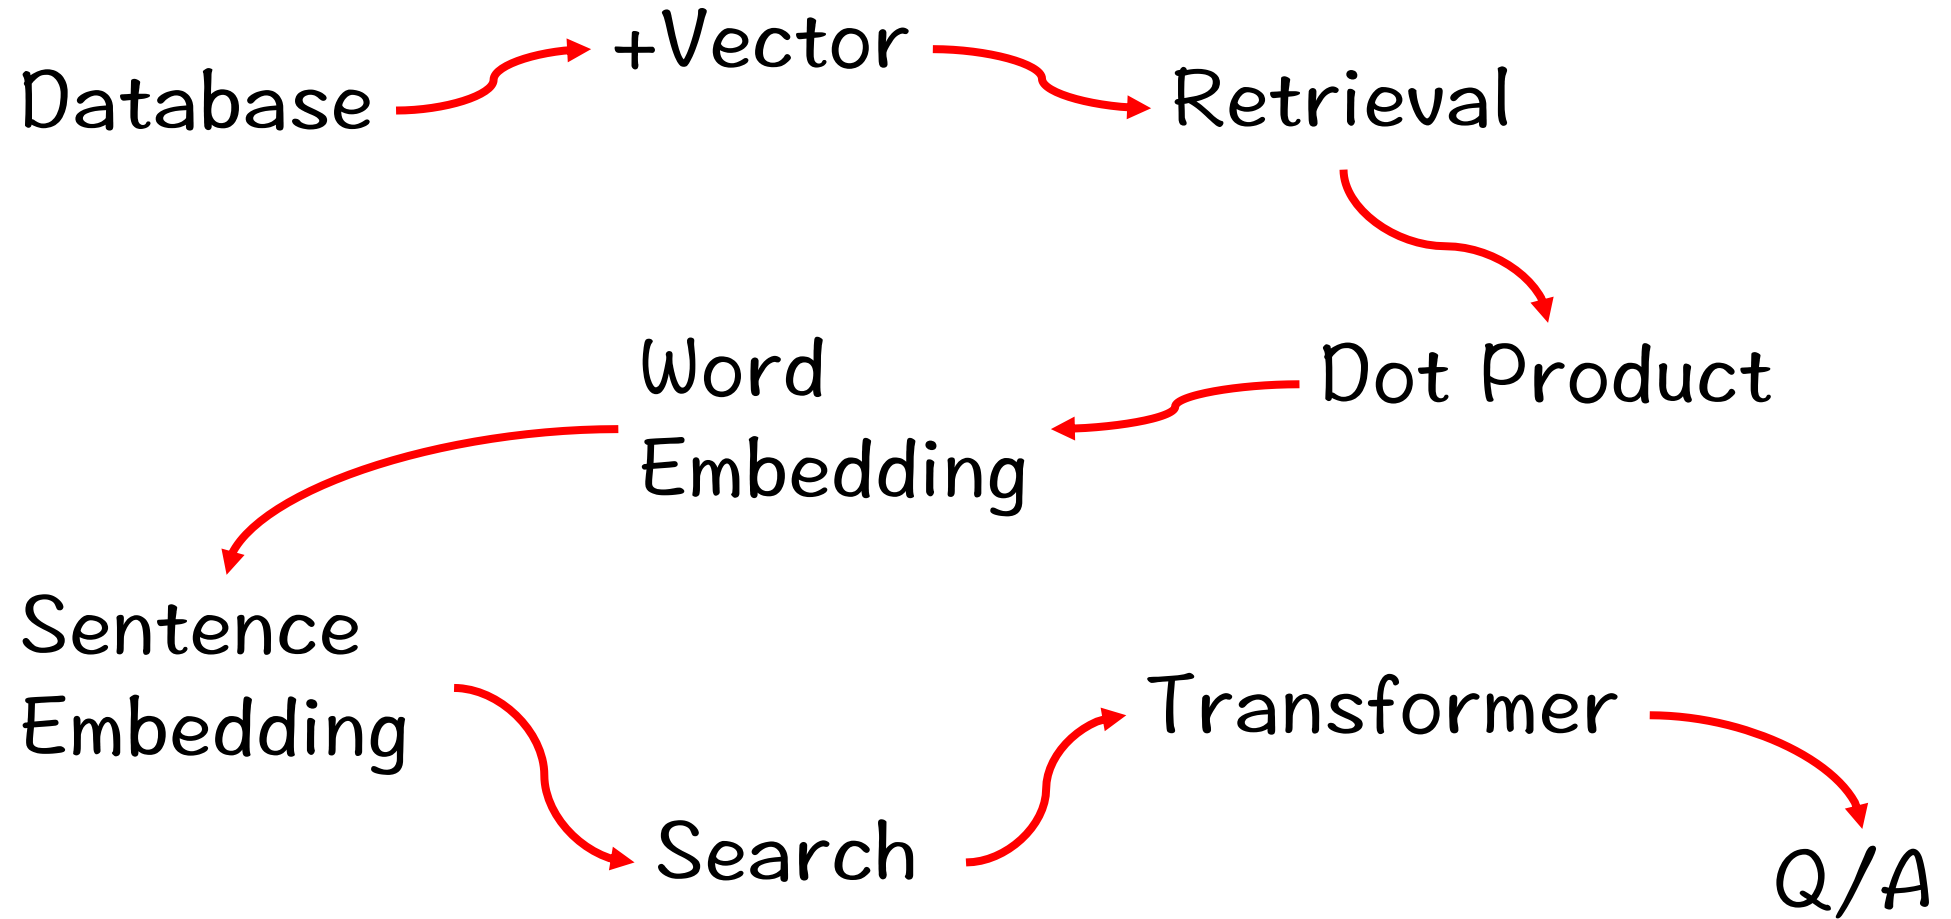



Beginner's Guide to Vector Databases

AI by Hand 🖋️

Prof. Tom Yeh

Roadmap





Database

Beginner's Guide to Vector Databases - AI by Hand 🖋️



University of Colorado
Boulder

Fun fact

There are _____ millions dogs in the world!

How to create a table?

SQL:

```
_____TABLE_____  
( id _____,  
  name _____,  
  size _____,  
  pop _____)
```


id	name	size	pop

How to insert a record?

SQL:

_____ INTO animals
_____ (1, dog, 2, 900)

id	name	size	pop
1	dog	2	900



Vector Database

Beginner's Guide to Vector Databases - AI by Hand 🖋️



University of Colorado
Boulder

How to create a vector database?

SQL: `CREATE TABLE animals`
`(id INT,`
`name VARCHAR(10),`
`size INT,`
`pop INT,`
`emb _____ not null)`

id	name	size	pop	


How to insert a record with a vector?

SQL:

INSERT INTO animals

VALUES (1, dog, 2, 900, _____)

id	name	size	pop	emb			
1	dog	2	900	<table><tr><td></td><td></td><td></td></tr></table>			



Retrieval

Beginner's Guide to Vector Databases - AI by Hand 🖋️



University of Colorado
Boulder

Which record is relevant to the query “cat”?


Query

cat			
1	2	0	


id	name	size	pop	emb			
1	dog	2	900	<table><tr><td>2</td><td>1</td><td>0</td></tr></table>	2	1	0
2	1	0					
2	bat	1	10000	<table><tr><td>0</td><td>1</td><td>2</td></tr></table>	0	1	2
0	1	2					

Draw distance vs similarity


distance



similarity



Distance vs similarity on a scale of 1 to 5



How to retrieve by similarity? (dot product)


_____ name, emb<____>[____,____,____] AS score

FROM animals

_____ BY _____ ASC | DESC ;|

How to retrieve by distance? (Euclidean)

```
SELECT name, emb<*>[1, 2, 0] AS score  
FROM animals  
ORDER BY score DESC;
```



Dot Product

Beginner's Guide to Vector Databases - AI by Hand 🖋️



University of Colorado
Boulder

How to compute dot product?

Example:

1	2	3
---	---	---

* * *

2	2	0
---	---	---

= = = Σ

2	4	0
---	---	---

6

Result

dog

2	1	0
---	---	---

* * *

cat

1	2	0
---	---	---

= = = Σ

--	--	--

Result

How to compute dot product using matrix multiplication?

Example:

			1
			2
			3
2	2	0	6


				dog
			2	
			1	
			0	
cat	1	2	0	

How to compute dot product with multiple vectors?

Example:

	1	1
	2	1
	3	1
2	2	0
6	4	

	dog	bat
	2	0
	1	1
	0	2
cat	1	2
	4	




Word Embedding

Beginner's Guide to Vector Databases - AI by Hand 🖋️




University of Colorado
Boulder

Where are dog, cat and bat in the “name” space?



Where are dog, cat and bat in the “name” space?



Which embedding is better?

Embedding 1

dog	cat	bat
2	1	0
1	2	1
0	0	2

Embedding 2

dog	cat	bat
2	0	1
1	1	0
0	2	2

Which embedding is better?

Desired
dot
product
similarity

dog	cat	bat
2	1	0
1	2	1
0	0	2

Embedding 1

dog	cat	bat
2	1	0
1	2	1
0	0	2

Embedding 2


dog	cat	bat
2	0	1
1	1	0
0	2	2

dog	2	1	0
cat	1	2	0
bat	0	1	2

	H	L
H		L
L	L	

dog	2	1	0
cat	1	2	0
bat	0	1	2

dog	2	1	0
cat	0	1	2
bat	1	0	2



Sentence Embedding

Beginner's Guide to Vector Databases - AI by Hand 🖋️



University of Colorado
Boulder

How to embed sentences?

id	comment	user	emb
1	How are you?	John	?
2	Who are you?	Mary	?

“How are you” → word embedding vectors

how	are	you
-----	-----	-----

a	an	the	how	why	who	what	are	is	am	be	was	you	we	I	they	she	he	she	me	him	her
0	-1	0	1	0	1	0	0	-1	1	0	0	0	3	1	0	-1	0	0	0	-1	0
2	0	2	0	0	0	-1	1	0	0	0	2	1	0	2	0	2	0	0	2	0	0
-1	0	-1	1	2	0	0	1	0	1	-1	0	0	-1	0	3	0	0	-1	0	2	-1
0	1	0	0	1	0	1	0	1	0	1	-2	0	0	0	1	0	1	0	1	0	1

Word vectors \rightarrow Sentence vector

Method 1: Concatenate

how	are	you
-----	-----	-----

1	0	0
0	1	1
1	1	0
0	0	0

Word vectors → Sentence vector

Method 2: Average

how	are	you
1	0	0
0	1	1
1	1	0
0	0	0

id	comment	user	emb
1	How are you?	John	
2	Who are you?	Mary	

“Who are you” → word embedding vectors

who	are	you
-----	-----	-----

a	an	the	how	why	who	what	are	is	am	be	was	you	we	I	they	she	he	she	me	him	her
0	-1	0	1	0	1	0	0	-1	1	0	0	0	3	1	0	-1	0	0	0	-1	0
2	0	2	0	0	0	-1	1	0	0	0	2	1	0	2	0	2	0	0	2	0	0
-1	0	-1	1	2	0	0	1	0	1	-1	0	0	-1	0	3	0	0	-1	0	2	-1
0	1	0	0	1	0	1	0	1	0	1	-2	0	0	0	1	0	1	0	1	0	1

	0	0
	1	1
	1	0
	0	0

Word vectors \rightarrow Sentence vector

Method 2: Average

who	are	you
1	0	0
0	1	1
0	1	0
0	0	0

id	comment	user	emb
1	How are you?	John	$[1/3, 2/3, 2/3, 0]$
2	Who are you?	Mary	

How to query by SQL?


```
_____ comment, emb<____>[____,____,____,____] AS score  
  
FROM posts  
  
ORDER BY _____ ASC | DESC ;
```


How to query using a high-level API?

```
query = Query(post_index)
        ._____ (post)
        ._____ (relevance_space.text, Param("_____"))

app.query(query, _____ = "who are you?" )
```

Source: Superlinked.com 



Search

Beginner's Guide to Vector Databases - AI by Hand 🖋️



University of Colorado
Boulder

K-Nearest Neighbor, K=3, Dot-Product

Database

ID	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20
emb																				

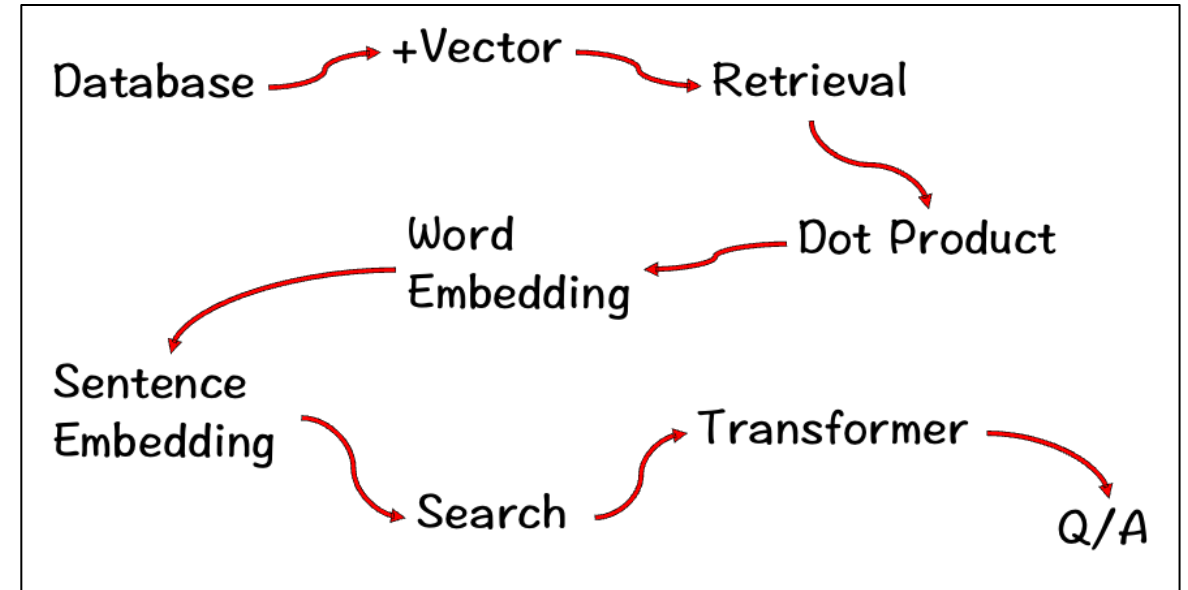
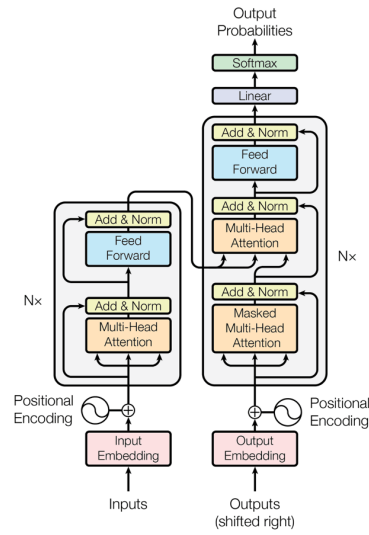
Query					9	-8	9	9	0	3	1	-6	0	11	3	13	-2	6	15	-9	7	6	-5	8
	{ max min }																							

K-Nearest Neighbor, $K=3$, L2

Database

ID	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20
emb																				

Query					6	8	9	9	1	10	0	9	12	15	2	13	12	6	15	9	7	6	5	8
{ max min }																								



Transformer

Beginner's Guide to Vector Databases - AI by Hand 🖋️

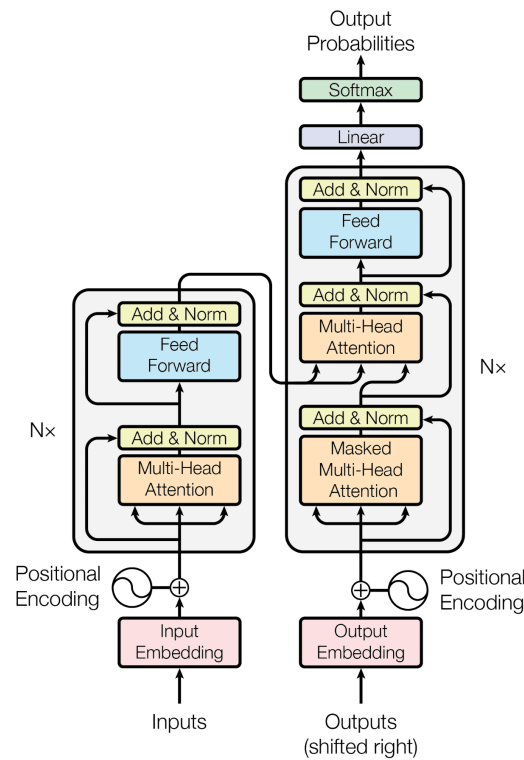


University of Colorado
Boulder

How to use a Transformer to get a sentence embedding vector?

Word
Embedding
Vectors

1	0	0
0	1	1
1	1	0
0	0	0



Sentence
Embedding
Vector



How to combine across positions?

1	0	0	1
0	1	1	0
1	1	0	1
0	0	0	

How to combine across positions?

			1	0
			0	1
			1	1
1	0	0	1	
0	1	1	1	
1	1	0	1	
0	0	0	0	

How to combine across positions?

			1	0	0
			0	1	0
			1	1	1
1	0	0	1	0	
0	1	1	1	2	
1	1	0	1	1	
0	0	0	0	0	

How to combine across features?

1	0	-1	0	1
1	0	0		
0	1	1		
1	1	0		
0	0	0		
1	1	1		

How to combine across features?

1	0	-1	0	1
0	1	1	0	0

1	0	0
0	1	1
1	1	0
0	0	0
1	1	1

1	0	1

How to combine across positions and features?

					1	0	0
					0	1	0
					1	1	1
	1	0	0			0	0
	0	1	1			2	1
	1	1	0			1	0
	0	0	0			0	0
					1	1	1
1	0	-1	0	1			

How to use a Transformer to get a sentence embedding vector?

Word
Embedding
Vectors

1	0	0
0	1	0
1	1	1

1	0	0
0	1	1
1	1	0
0	0	0

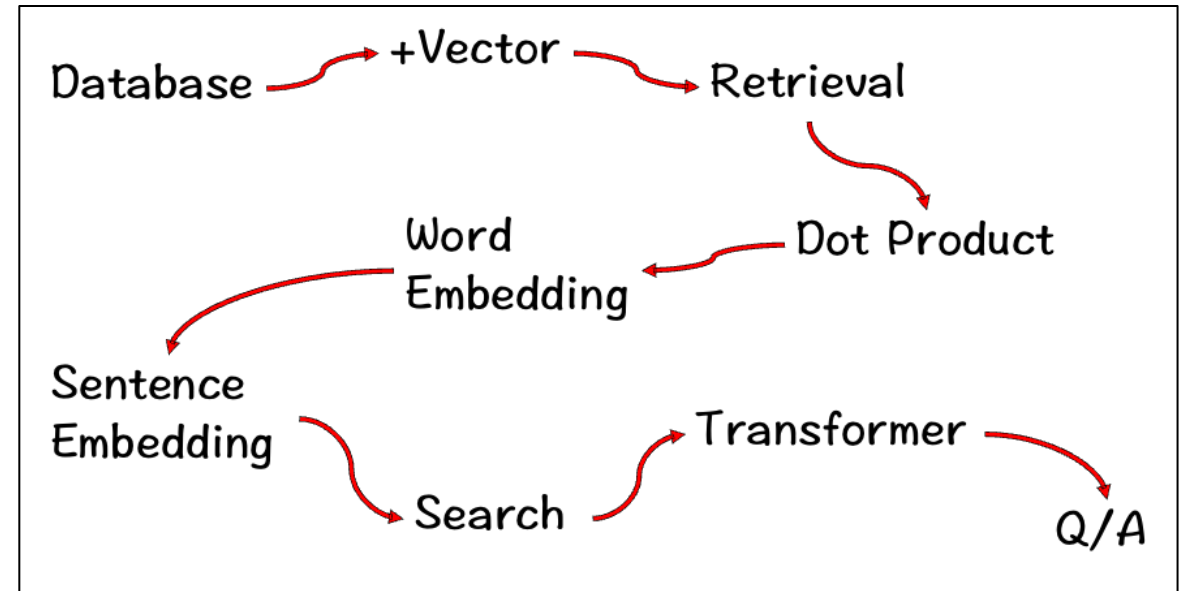
	0	0
	2	1
	1	0
	0	0

1 1 1

Sentence
Embedding
Vector

1	0	-1	0	1
0	1	1	0	0
0	0	0	1	1
0	0	1	1	0

2	3	1
1	1	1
1	1	0



Q/A

Beginner's Guide to Vector Databases - AI by Hand 🖋️



University of Colorado
Boulder

